

Real Time Embedded Systems

Worksheet 4. The Time-Slicing Structure

This week we start work on the central components of an elementary real-time operating system - we will call it a 'runtime system' - that divides the processor's time between separate user tasks. These tasks will need to communicate with the system, and will request its services by means of a 'software interrupt'.

Implementation of Software Interrupts

A software interrupt is known as a 'trap'. It causes the processor to respond in a very similar way as it does to a hardware interrupt, and so allows system calls from the user program and interrupts from hardware devices to enter the operating system in a consistent way.

There are 16 trap instructions available, numbered 0 to 15, and written

```
trap #0
...
trap #15
```

Each of the 16 trap instructions may have its own interrupt service routine (ISR). After pushing the PC and SR, the processor then accesses a table in low memory, at address 80H. As for the hardware interrupts, the table contains a 4-byte value corresponding to the address of the ISR for each software interrupt. Vectors for the two types of interrupt will normally be combined into a single block of code.

```

;interrupt vectors
org    $64    ;origin 64H
hvec1  dc.l   hisr1 ;address of hardware ISR 1
hvec2  dc.l   hisr2 ; ... etc

org    $80    ;origin 80H
svec0  dc.l   sivr0 ;address of software ISR 0
svec1  dc.l   sivr1 ; ... etc
```

Controlling Interrupts

There is, however, an important difference between hardware and software interrupts. Hardware interrupts are in order of priority, with 7 being the highest priority and 1 the lowest. If two hardware interrupts occur at the same time, then the one at the higher priority will be accepted and the other one will be kept waiting until the first ISR has completed. If a hardware interrupt occurs shortly after another one, but while the ISR for the first interrupt is still in execution, then the processor will again compare the priorities of the two interrupts. If the new interrupt is of a higher priority, then it will interrupt the lower priority ISR. If the new interrupt is at a lower priority than the currently executing ISR, it will be kept waiting until that ISR completes.

Software interrupts do not behave in an analogous way. Since the processor can only execute one instruction at a time, it would be impossible for two software interrupts to occur at the same time, and unless a programmer includes a trap instruction within an ISR, there will also be no occasions on which a trap takes place during the processing of another trap. There is therefore no point in prioritising the software interrupts, and all 16 are at the same priority. There is, however, the question of the relative priority of the hardware and software interrupts. What if a hardware interrupt is raised at the same time as the processor is executing a software interrupt instruction? This is

handled by assigning all the software interrupts to priority level 0. Processing of a software interrupt is therefore interruptible by a hardware interrupt at any of the priority levels 1 to 7.

Within your system, however, regardless of the type of interrupt being processed, you will want to prevent the acceptance of any other interrupt. Your system will therefore be completely un interruptible. Once entered, it will always run to completion and then return to the user task that was running when the interrupt was raised. You will therefore need to disable interrupt acceptance, the procedure for which is explained now.

Using the simulator, examine the 16-bit status register. Bits 8, 9 and 10 (labelled 'INT') hold a 3-bit value that represents the interrupt priority mask. When an interrupt is accepted, the mask is set to the priority level of that interrupt. A hardware interrupt will only be accepted if its priority is greater than the current setting in the mask. Normally, the mask is set to 000 (decimal 0) thereby allowing the acceptance of any hardware interrupt. However, it will remain at zero during its response to a software interrupt, since that is the priority of these interrupts, and will thereby allow the hardware to interrupt the software ISR. If you wish to prevent this, then the following instruction, placed at the very start of a software ISR, sets the mask to binary 111 (decimal 7). Any hardware interrupts will now be disabled, and held pending until the mask is returned to zero.

```
or    #$0700,sr    ;disable hardware interrupts
```

The status register will have been automatically saved on the stack at the start of the interrupt servicing. On execution of the 'return from exception' instruction (RTE), it will be restored, and the mask reset to the zero value that it held previously, thereby allowing the acceptance of any hardware interrupt that might have been raised in the meantime and is currently pending.

If you want to enable hardware interrupts at any other time, the following instruction will set the mask to zero.

```
and    #$f8ff,sr    ;enable hardware interrupts
```

Add WeChat powcoder
Practical Work

Assessment question

*Work in pairs on this question, and keep a copy of your answer. You will need to submit your software, including the test programmes that you use to demonstrate it, and your documentation. These items should be placed into a single zipped file, and uploaded to a Canvas submission point to be advised. The submission deadline is **2pm on Friday 20th January, 2023**.*

The work consists of writing a basic time-slicing system, along the lines of the one discussed in the lecture. It should allow the execution of several concurrent user tasks, with support for task scheduling and inter-task communication. Test your system, and using short test programmes of the type used in the lectures, show clearly that each function is working. Demonstrate the operation of mutual exclusion and task synchronisation. Write a short document describing how to use the system, stating the maximum number of instructions executed by each of the functions.

The system runs in the foreground, and is entered following either a timer interrupt or a software interrupt from one of the tasks requesting service. Another hardware interrupt will be added later.

The following system calls should be supported by means of software interrupts. They can either each be allocated to a separate trap number, or (as in the demonstration system) they can all be called on the same trap, with one of the registers used to hold a value identifying the requested function. Some of the calls also require additional parameters in other registers.

1. Create task

Function: A currently unused TCB is marked as in use and set up for a new task. It is placed on the ready list. The requesting task remains on the ready list.
Parameters: The start address of the new task,
The address of its top-of-stack.

2. Delete task

Function: The requesting task is terminated, its TCB is removed from the list and marked as unused.
Parameters: None.

3. Wait mutex

Function: If the mutex variable is one, it is set to zero and the requesting task is placed back onto the ready list. If the mutex is zero, the task is placed onto the wait list, and subsequently transferred back to the ready list when another task executes a *signal mutex*.
Parameters: None.

4. Signal mutex

Function: If the mutex variable is zero, and a task is waiting on the mutex, then that task is transferred to the ready list and the mutex remains at zero. If the mutex is zero and no task is waiting, the mutex is set to one. In either case, the requesting task remains on the ready list.
Parameters: None.

5. Initialise mutex

Function: The mutex is set to the value 0 or 1, as specified in the parameter.
Parameters: 0 or 1.

When this is working, extend the system with the following functions. Function 6 will require the use of an additional interrupt at level 2, and you will need to consider the implications of this for the correct working of the system.

6. Wait I/O:

Function: The requesting task is placed onto the wait list, until an interrupt signifies completion of an I/O operation, when the task is transferred back to the ready list.
Parameters: None.

7. Wait time

Function: The requesting task is placed onto the wait list until the passage of the number of timer interrupts specified in the parameter, when it is transferred back to the ready list.
Parameters: Number of timer intervals to wait.

An additional function is executed automatically at start-up, or if the user presses the reset button.

System reset

Function: The system is initialised: all internal variables are reset, and each TCB is marked as unused. A TCB for task T0 is then created, and T0 becomes the running task.

The system assumes that a default user task, T0, is present. The system runs this task immediately after a reset. It will need to be located at a predetermined address, which will be coded into the reset function.

Your system should be robust, and deal with errors in an intelligent way. For example, what if the user tries to create more tasks than there are available TCBs?

The Demonstration System

A system was demonstrated during the lecture. Excerpts of code from it are reproduced below. These can be used in your own systems; unchanged, modified, or rewritten as you wish.

The system recognises a hardware interrupt at level 1 from the timer. It also allows system calls by means of software interrupts, all of which have been allocated to trap 0. These system calls are programmed by placing a value that identifies the requested function into data register 0, and any other parameters as required by each of the individual functions in registers D1 onwards. For example, system call 1 is used to create a new task. Suppose that this new task is called T1, and that its top-of-stack is to be located at address 6000H. It would be programmed as follows.

```


move.l    #1,d0      ;set id in d0
move.l    #t1,d1     ;set address of new task in d1
move.l    #$6000,d2  ;set stack address in d2
trap      #0         ;call system

```

The main data structure used in this system is a list of task control blocks (TCBs). Each TCB represents the state of one of the tasks. It contains a copy of all that task's registers, together with some items of control information including a flag that indicates whether the TCB is in use.

At any time, each of the tasks will be in one of three states: it will be the currently running task, it will be ready to run when its turn comes up, or it will be unable to run because it is waiting for the occurrence of some event, which could be a signal operation on a mutex, the expiry of a time interval, or an I/O interrupt.

At initialisation, all the TCBs in the list are marked as unused. As each new task is started, one of the unused TCBs is allocated for it and marked as used. These TCBs are organised into two linked lists, in which each element contains a pointer to the next element. These lists are called 'ready' and 'waiting'. Two more data items consist of pointers to the first element in each list. The pointer `rdytcbl` holds the address of the first element in the list of ready TCBs. The first element in this list is the task that actually is running. The linkage in this list is circular, that is, the last entry points back to the first, so making it easy to access each TCB in rotation. The pointer `wttcbl` holds the address of the first element in the list of TCBs that are waiting. There is no need to access elements of this list in rotation, so the last element in this list has its pointer set to zero.



The example above shows a list of 8 elements, each of which has a pointer, labelled N, to the next element. Elements 4, 6, 0 and 1 are on the ready list, with element 4 being the TCB for the running task. Elements 3 and 7 are on the waiting list, and elements 2 and 5 are unused.

The system is organised into the following sections.

- Data definitions and equates
 - org 0
- Interrupt vectors
- Executable code
 - System reset
 - First-level interrupt handler
 - Service routines
 - Scheduler
 - Dispatcher
- Data storage
- Default user task T0

Data definitions

Each TCB represents the state of one of the current tasks, and is defined as follows.

```

tcb      org      0                      ;tcb record
tcdb0    ds.1     1                      ; D register save
tcdb1    ds.1     1
tcdb2    ds.1     1
tcdb3    ds.1     1
tcdb4    ds.1     1
tcdb5    ds.1     1
tcdb6    ds.1     1
tcdb7    ds.1     1
tcba0    ds.1     1                      ; A register save
tcba1    ds.1     1
tcba2    ds.1     1
tcba3    ds.1     1
tcba4    ds.1     1
tcba5    ds.1     1
tcba6    ds.1     1
tcba7    ds.1     1
tcbsr    ds.1     1                      ; SR (status reg) save
tcbps    ds.1     1                      ; PC save
tcnext   ds.1     1                      ; link to next record
tcused   ds.1     1                      ; record in use flag
          ds.1     1                      ; other fields as required
          ds.1     1                      ;
tcblen   equ      *                      ; length of tcb record in bytes

```

Data storage

Storage for a list of TCBs is defined as in the first line below. The constant `ntcb` represents the number of TCBs in the list, and should be set up as an equate. Other variables are described throughout these notes.

```

tcblist  ds.b      tcblen*ntcb           ;tcb list (length x no of tcbs)
rdytcb   ds.1      1                     ;^ ready tcb list
wt tcb   ds.1      1                     ;^ waiting tcb list
a0sav    ds.1      1                     ;A0 temporary save
d0sav    ds.1      1                     ;D0 temporary save
id        ds.1      1                     ;function id

```

Interrupt vectors

The interrupt vectors are addresses of the code that will be executed as a result of an interrupt. The following three addresses are defined.

Address `res` is the location of the routine to which the processor branches when the it responds to a hardware reset. Address `fltint` is the location to which the processor branches following a timer interrupt at level 1, and `flsint` following a software interrupt. The address `stk` is the value that is loaded into the stack pointer following a hardware reset.

```
; *****
;                                     ; INTERRUPT VECTORS
; *****

    org      0

    dc.l     stk                ; initial SP
    dc.l     res                ; reset
    org      $64
    dc.l     fltint             ; interrupt 1 (timer)
    org      $80
    dc.l     flsint             ; trap 0 (system call)
```

Executable Code

Assignment Project Exam Help

First-level interrupt handler

The first-level interrupt handler (FLIH) contains the code that is executed immediately following an interrupt. Hardware interrupts at level 1 are directed by the interrupt vector to enter the FLIH at `fltint`, while software interrupts arrive at `flsint`. The FLIH performs three main functions.

It takes the pointer to the TCB of the currently executing task, stored at `id_tcb`, and saves the values of the registers, including the PC and SR, within that TCB.

It also sets a value within a storage location, known as `id`, that identifies the source of the interrupt. If an interrupt has been raised by the hardware timer, then `id` is set to 0. For a software interrupt, `id` is set to the value, from 1 onwards, of the system call function number. The `id` will subsequently be used to select the corresponding service routine for processing this interrupt.

Programming the above two operations requires particular care, because saving the value of the user's registers as they were at the time of the interrupt requires the use of certain registers itself. Registers D0 and A0 are in use for this purpose. These registers are therefore saved in temporary locations, before being transferred to their long-term holding locations within the TCB.

The other function performed by the FLIH is to disable interrupts, if this has not already happened. A level-1 hardware interrupt from the timer will have set the interrupt priority mask to 1, thereby preventing any further interrupts. A software interrupt will have left the mask at 0, which would allow the timer device to interrupt the processing of the software interrupt. Therefore the first action taken at the software interrupt entry point is to disable hardware interrupts by setting the mask to 7.

```

;*****
flih                                ;FIRST-LEVEL INTERRUPT HANDLER
;*****

fltint                                ;ENTRY FROM TIMER INTERRUPT
    move.l    d0,d0sav                ;save D0
    move.l    #$0,d0                  ;set id = 0
    move.l    d0,id
    move.l    d0sav,d0                ;restore D0
    bra       fl1

flsint                                ;ENTRY FROM TRAP (SOFTWARE INTERRUPT)
    or        #%0000011100000000,sr ;disable hardware interrupts
    move.l    d0,id                  ;store id
    bra       fl1

fl1
    move.l    a0,a0sav                ;save working reg

    move.l    rdytcb,a0              ;A0 ^ 1st ready tcb (ie running tcb)

    move.l    d0,tcbd0(a0)            ;store registers
    move.l    d1,tcbd1(a0)
    move.l    d2,tcbd2(a0)
    move.l    d3,tcbd3(a0)
    move.l    d4,tcbd4(a0)
    move.l    d5,tcbd5(a0)
    move.l    d6,tcbd6(a0)
    move.l    d7,tcbd7(a0)
    move.l    a0sav,d0
    move.l    d0,tcba0(a0)
    move.l    a1,tcba1(a0)
    move.l    a2,tcba2(a0)
    move.l    a3,tcba3(a0)
    move.l    a4,tcba4(a0)
    move.l    a5,tcba5(a0)
    move.l    a6,tcba6(a0)

    move       (sp),d0                ;pop and store SR
    add.l      #2,sp
    move.l     d0,tcbstr(a0)

    move.l     (sp),d0                ;pop and store PC
    add.l      #4,sp
    move.l     d0,tcbpc(a0)

    move.l     a7,tcba7(a0)           ;store SP

;START OF SERVICE ROUTINES

```

System reset and service routines

The service routines are arranged as a large switch statement, using `id` as the case variable. Each routine carries out one of the functions defined in the specification.

Scheduler

The scheduler examines the ready list, to which `rdytcb` points to the first element. This is the TCB of the task that was executing when the system was invoked, and which has just been interrupted. By following the links, the scheduler can locate each TCB that is currently ready to run. It selects

one of these tasks for running, and adjusts the value in `rdytc_b` to point to the TCB for this task. This TCB will be then used by the dispatcher to resume execution of the task.

The scheduler may make the decision as to which task will run next by doing nothing more than following the link in the current TCB to the next one in the chain. This will result in each ready task running in rotation, receiving an approximately equal amount of run time each. Alternatively, it would be possible to assign a priority to each task as it is created, by adding another parameter to the 'create task' system call. Higher priority tasks would then receive a larger proportion of the available run time.

Dispatcher

The dispatcher reverses the action taken by the FLIH. Using the newly set value in `rdytc_b`, it restores the registers of the selected task to the values that were stored when that task was interrupted. Careful housekeeping is again necessary, as this operation itself requires the use of registers D0 and A0. The dispatcher finishes by recreating the state of the stack as it was after the task was interrupted. The processor then uses a 'return from exception' instruction, as though it were returning from any normal interrupt, to transfer control back to the selected task.

```
;
;END OF SCHEDULER

;*****
disp
;*****

move.l rdytc_b,a0          ;A0 ^ new running tcb
move.l tcba1(a0),d1        ;restore registers
move.l tcba2(a0),d2
move.l tcba3(a0),d3
move.l tcba4(a0),d4
move.l tcba5(a0),d5
move.l tcba6(a0),d6
move.l tcba7(a0),d7
move.l tcba1(a0),a1
move.l tcba2(a0),a2
move.l tcba3(a0),a3
move.l tcba4(a0),a4
move.l tcba5(a0),a5
move.l tcba6(a0),a6
move.l tcba7(a0),a7

sub.l #4,sp                ;push PC
move.l tcbpc(a0),d0
move.l d0,(sp)

sub.l #2,sp
move.l tcbstr(a0),d0       ;push SR
move.l d0,(sp)

move.l tcba0(a0),d0        ;restore remaining registers
move.l tcba0(a0),a0

rte                        ;return
```


An example of a user programme running under this system is shown here. It consists of two concurrent tasks. Task T0 calls the system to start task T1, then switches on the RH LED. Task T1 calls the system to wait for 4 timer intervals, then switches on the LH LED. From then on, the two tasks run alternately. If the timer is set to interrupt at one-second intervals, the result is that the RH LED lights immediately, then after 4 seconds the two LEDs start alternating.

```

;*****
;                                     ;USER APPLICATION TASKS
;*****

                                ;system call equates
sys      equ      0            ; system call trap (trap 0)
sysst    equ      1            ; start new task
system   equ      2            ; terminate task
syswtim  equ      3            ; wait on timer

;*****
;                                     ;USER APPLICATION TASKS
;*****

        org      usrcode

led      equ      $e00010      ;led
sw       equ      $e00014      ;switch

t0:      ;TASK 0
        move.l   #sysst,d0      ;start task 1
        move.l   #t1,d1         ; address
        move.l   #$4000,d2      ; top of stack
        trap     #sys           ;repeat
t00:      move.l   #$01,d1        ; set led 0
        move.b   d1,led
        bra      t00

t1:      ;TASK 1
        move.l   #syswtim,d0     ;wait for 4 clocks
        move.l   #4,d1
        trap     #sys
t10:      ;repeat
        move.l   #$02,d0        ; set led 1
        move.b   d0,led
        bra      t10

        END      res

```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder